

---

# CSE151B Final Report: ML Noobs Group

---

Jeffrey Feng || Dennis Wu || Tat Hei Tsin || Yujian He  
A15507377 || A15537107 || A15541127 || A15386248

## Abstract

This is the mile stone report for CSE151B deep learning project for Autonomous vehicles(AV). We are given 19 timestamps information for all agent such as position and velocity and expected to predict the next 30 time steps. We include my approach using deep learning model and linear regression with supportive data.

## 1 Task Description and Background

**1.1 Describe in your own words what the deep learning task is and why it is important. Provide some real-world examples where solving this task can have great impact on our daily life and the society.**

The background of the task focuses on a spatial-temporal task, where the deep learning model is required to predict future trajectories of moving objects like cars on roads based on their previous positions and velocity. By improving the accuracy of the prediction, it could benefit the forecasting technology of fully autonomous driving vehicles, which would increase the safety of driving without supervision. Therefore, solving this problem can boost the progression of autonomous driving technology. Moreover, in the domain of traffic forecasting, predicting the trajectories of vehicles could give a better prediction and understanding of traffic conditions in major cities, thus benefiting efficiency of highway and public transportation planning.

**1.2 Use Google Scholar or other internet resources to research on this task. What type of methods have been examined before? Include some references and discuss their ideas in a few sentences. You can use Bibtex to manage your bibliography**

Across the methods that have been applied on this task before, convolutional neural networks (CNN) (Ma et al., 2017) (Yu et al., 2017), recurrent neural networks (RNN) (Cui et al, 2018), variants like long-short term memory model (LSTM) (Feng et al., 2019), autoencoder like Seq2seq, and graphical convolutional networks (GCN) (Zhao et al., 2017) are all applied on this task. We would focus on works that are more interesting to use, which includes the T-GCN by Lehai Feng that predicts urban traffic flow using Graph Convolutional Network (Feng et al., 2019), the work by Zhao et al. that predicts the short-term traffic using LSTM network (Zhao et al., 2017), the model built by Yu et al. that forecasts the geospatial coordinates of traffic using Convolutional Neural Network (Yu et al., 2017). Correspondingly, we thought of building our model based on CNN, RNN and Linear Regression. Feng's T-GCN transforms the dataset into a graph, which can be further learned by the CNN model that enables us to predict the car's trajectory. The work by Zhao et al. made us realize the power of the RNN by training the model using LSTM.

**1.3 Define the input and output in mathematical language and formulate your prediction task. From the abstraction, do you think your model can potentially solve other tasks beyond this project? If so, list some examples and explain your rationale**

For this task, we are given a dataset that consists of individual scenes. Within each scene, the input matrix for the model can be described as  $m \times n \times p$ , where  $m$  is the number of autonomous vehicles

being tracked within the scene.  $m = 60$  in all our cases, but there are cases where the number of tracked vehicles is smaller than 60.  $N$  equals the number of time steps recorded, where the recording interval is between 2 seconds. In all the cases,  $n = 19$ .  $p$  equals to the features the model is going to intake, which includes  $p_i n$  and  $v_i n$ .  $p_i n$  contains a vector  $[x_i, y_i]$  which records the vehicle's positions for  $x$  and  $y$  coordinates.  $v_i n$  contains a vector  $[v_x, v_y]$  which recorded the vehicle's position for  $x$  and  $y$  velocity. We are given the first 19 times steps of vehicles' information to predict the next 30 time steps of a corresponding target vehicle. The output matrix would have a similar form of  $(m \times n \times p)$ , where  $n = 30$ .

Other than prediction of the car coordinates, this deep learning model can solve any other tasks that record the position and velocity of objects along time, since this is what our deep learning model is designed for. For example, we can transform the electricity usage over a number of households, and predict the next 30 day dynamic electricity usage using our model. We can also predict fluctuations of the stock market or generate new sentences according to previous records.

## 2 Exploratory Data Analysis : Perform exploratory data analysis and report your findings with texts/tables/figures. If you include more exploratory analysis beyond the listed questions that provides insights into the data, you will receive bonus points.

### 2.1 Perform exploratory data analysis and report your findings with texts/tables/figures. If you include more exploratory analysis beyond the listed questions that provides insights into the data, you will receive bonus points.

#### 2.1.1 What is the train/test data size?

Train set dimension is  $205942 * 60 * 19 * 2$ , the test set dimension is  $3200 * 60 * 19 * 2$

	city	lane	lane_norm	scene_idx	agent_id	car_mask	p_in	v_in	track_id
0	scalar	(144, 3)	(144, 3)	scalar	scalar	(60, 1)	(60, 19, 2)	(60, 19, 2)	(60, 30, 1)
1	scalar	(810, 3)	(810, 3)	scalar	scalar	(60, 1)	(60, 19, 2)	(60, 19, 2)	(60, 30, 1)
2	scalar	(450, 3)	(450, 3)	scalar	scalar	(60, 1)	(60, 19, 2)	(60, 19, 2)	(60, 30, 1)
3	scalar	(252, 3)	(252, 3)	scalar	scalar	(60, 1)	(60, 19, 2)	(60, 19, 2)	(60, 30, 1)
4	scalar	(126, 3)	(126, 3)	scalar	scalar	(60, 1)	(60, 19, 2)	(60, 19, 2)	(60, 30, 1)
5	scalar	(99, 3)	(99, 3)	scalar	scalar	(60, 1)	(60, 19, 2)	(60, 19, 2)	(60, 30, 1)

This is the dimension mock-up of the first sixth scene.

#### 2.1.2 how many dimensions of inputs/outputs in the raw data?

inputs:  $60 * 19 * 2$   
outputs:  $60 * 30 * 2$

#### 2.1.3 what are the meanings of these input/output dimensions?

This means that for each scene, there are 60 cars in total (although some of the 60 cars will be empty), and for each of the car in the scene, there will be 19/30 frames of data depending whether it's input or output. Last but not least, for each frame of the cars, there will be 2 inputs in total,  $x$  and  $y$  coordinates.

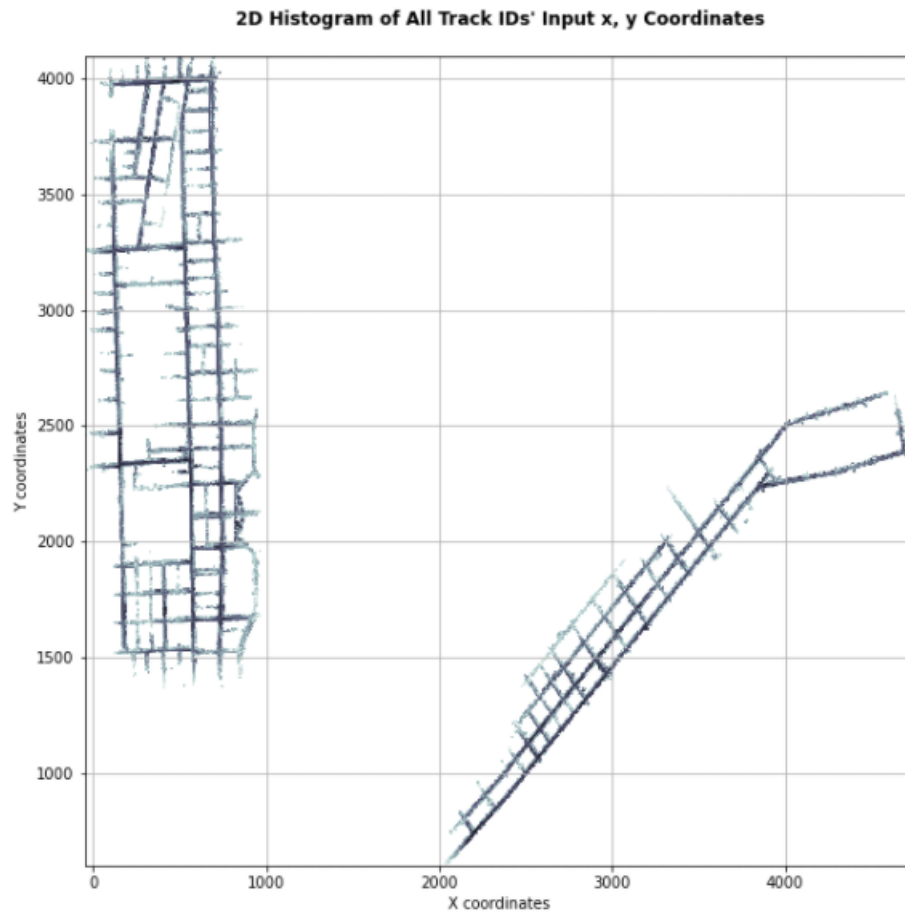
70 **2.1.4 what does one data sample looks like?**

```
{'city': 'MIA',
 'lane': array([[ 559.5046 , 4004.9236 ,    0.    ],
                [ 550.3973 , 4004.625  ,    0.    ],
                [ 541.29004, 4004.3262 ,    0.    ],
                [ 532.1871 , 4003.9185 ,    0.    ],
                [ 523.08466, 4003.4966 ,    0.    ],
                [ 546.20245, 3982.5588 ,    0.    ]....., dtype=float32),
 'lane_norm': array([[-9.10728359e+00, -2.98760027e-01,  0.00000000e+00],
                    [-9.10728359e+00, -2.98760027e-01,  0.00000000e+00],
                    [-9.10728359e+00, -2.98760027e-01,  0.00000000e+00],
                    [-9.10296726e+00, -4.07719672e-01,  0.00000000e+00],
                    [-9.10240650e+00, -4.21878159e-01,  0.00000000e+00],
                    [-9.10233021e+00, -4.23503131e-01,  0.00000000e+00],
                    [-7.13277912e+00,  3.98584557e+00,  0.00000000e+00]....., dtype=float32),
 'scene_idx': 44393,
 'agent_id': '00000000-0000-0000-0000-0000000032566',
 'car_mask': array([[1.],
                    [1.],
                    [1.]....., dtype=float32),
 'p_in': array([[ 501.16738892, 4017.24853516],
                [ 501.16748047, 4017.24853516],
                [ 501.16751099, 4017.2487793 ],
                ...]),
 'v_in': array([[ 3.13575329e-05,  1.31237772e-04],
                [ 7.79712456e-04, -4.57685994e-04],
                [ 3.82767001e-04,  1.66631851e-03],
                ...]),
 'p_out': array([[ 501.16738892, 4017.24902344],
                [ 501.16744995, 4017.24902344],
                [ 501.16760254, 4017.24926758],
                ...]),
 'v_out': array([[-2.96797277e-03,  1.41289853e-03],
                [ 5.01174422e-04,  8.26028525e-04],
                [ 1.55860023e-03,  3.85732530e-03],
                ...]),
 'track_id': array([[['00000000-0000-0000-0000-000000000000'],
                    ['00000000-0000-0000-0000-000000000000'],
                    ['00000000-0000-0000-0000-000000000000'],
                    ...], dtype=object])}
```

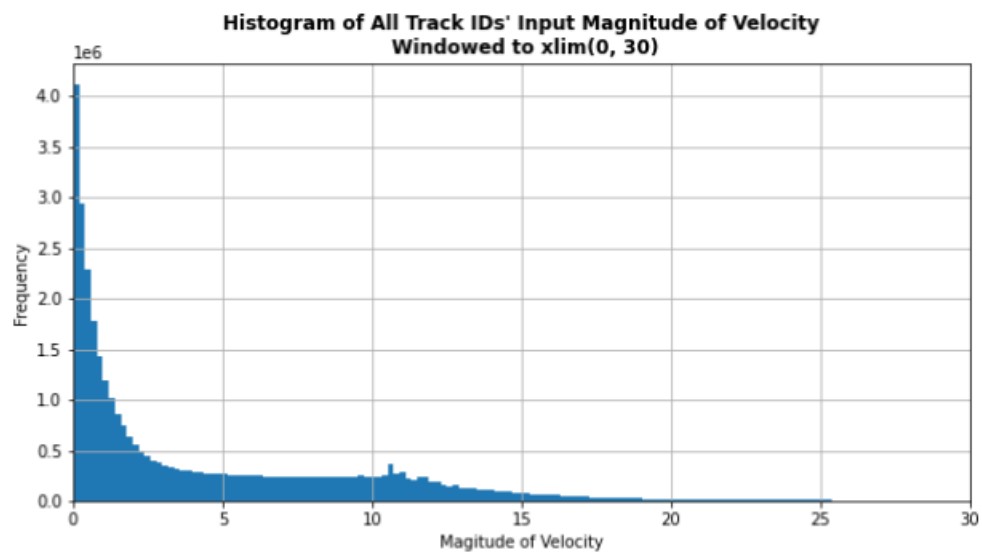
71

72 **2.2** statistical analysis to understand the properties of the data.

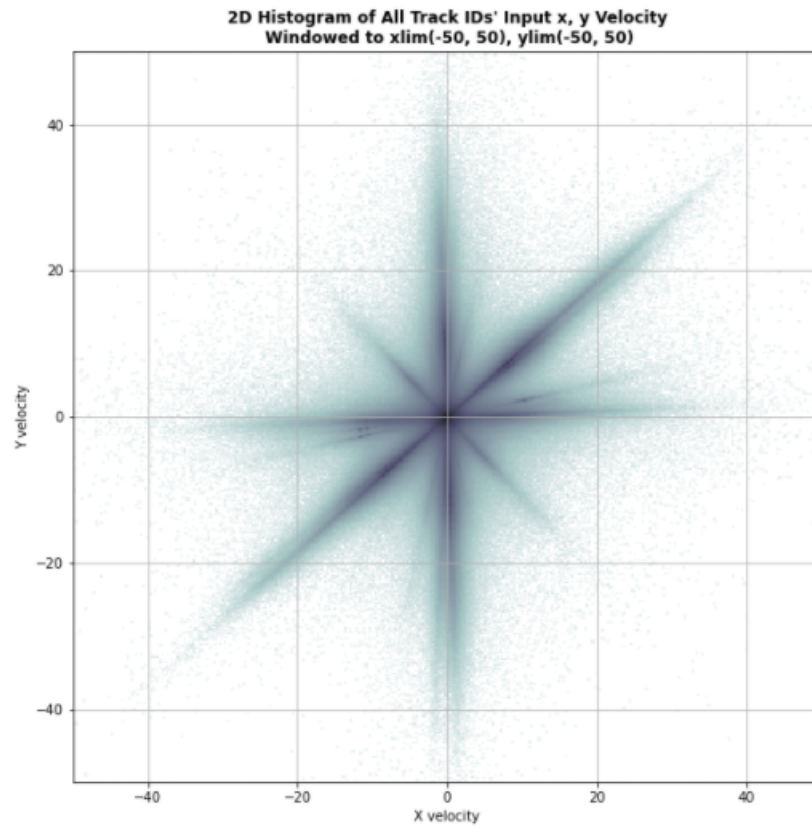
73 **2.2.1** what is the distribution of input positions/velocity (magnitude) for all agents?



74

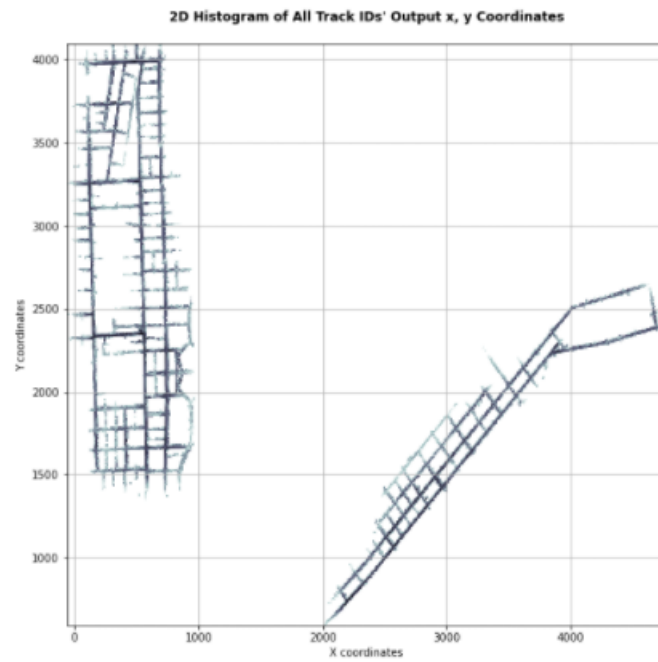


75



76

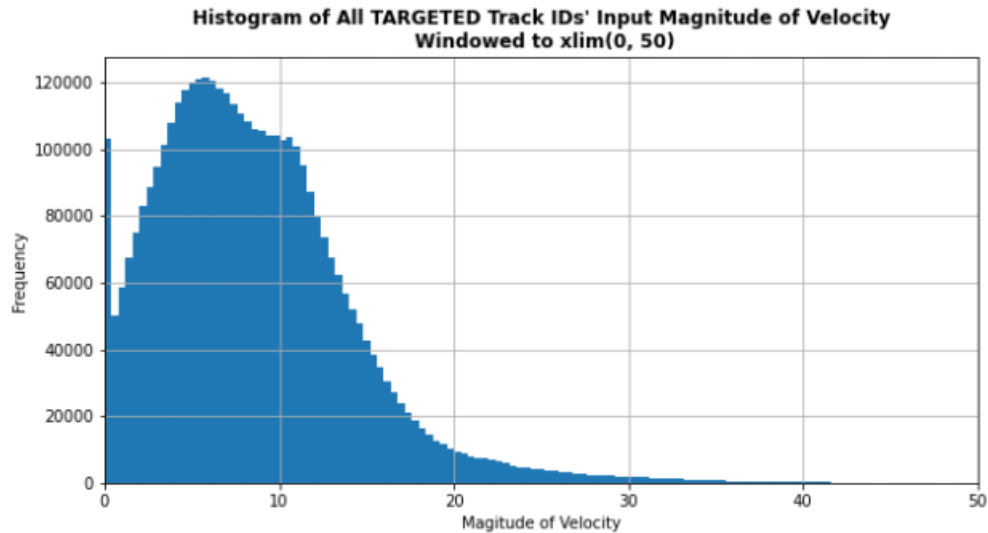
77 **2.2.2** what is the distribution of output positions/velocity (magnitude) for all agents?



78

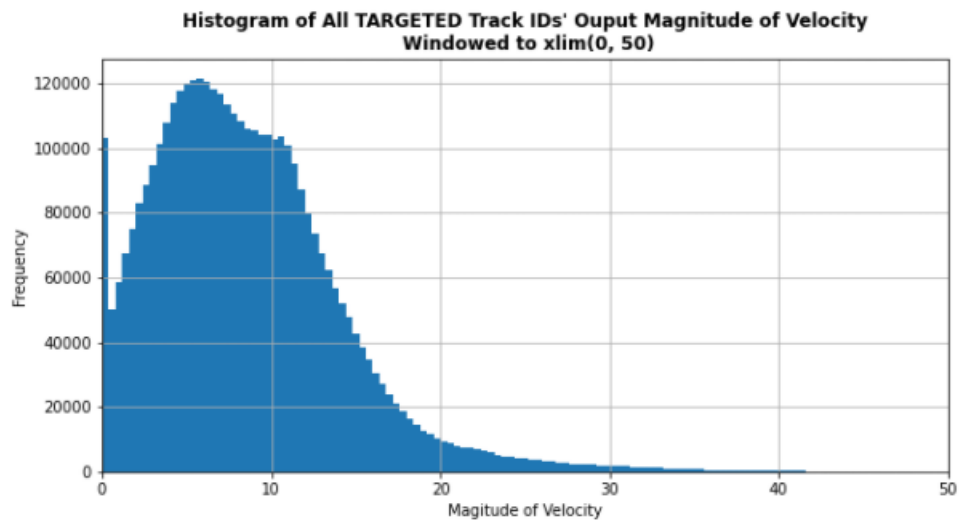
79 **2.2.3 what is the distribution of positions/velocity (magnitude) for the target agent?**

80



81

82



83

84 **2.3 Process the data to prepare for the prediction task. Describe the steps that you have**  
85 **taken to process the data. Your description should at least answer the following**  
86 **questions**

87 **2.3.1 Did you use any feature engineering? If yes, how did you design your features?**  
88 **Explain your rationale**

89 For our baseline model, in order to increase the model efficiency, we used linear regression and  
90 dropped all the car where car-mask =0, This lead to 181217 instances and we are able to split into  
91 small batches using dataloader class in pyTorch. According to physics, and our validation, the x,  
92 y coordinates along should cover the weight of the velocity, which should be able to estimated by  
93 using distance difference over time, so we did not use the velocity variables; Other variables, such  
94 as the lane information, were not included in the model (the dimensions for every lane is different)  
95 The only variables we used are the x and y locations. However, instead of the normalized location

96 value or the raw location values, we find the difference between the current location and the previous  
97 location, resulting in  $(19 - 1)$  difference in x values, and  $(19 - 1)$  difference in y values; We did not  
98 drop `car_mask = 0` for the later RNN model because it provides easier dimension manipulation.

### 99 **2.3.2 How did you normalize your data? Why did you choose this normalization scheme?**

100 We tried normalization of the x and y coordinates by using the mutual mean between the both cities,  
101 but after thinking about the overlapping of the features on the map, we decided not to use the mutual  
102 mean. Instead, we normalized the coordinates based on which cities the scene is in. However, for the  
103 linear model that yields the best result, we took the difference between the current location and the  
104 previous location, resulting in  $(19 - 1)$  entries in x locations, and  $(19 - 1)$  entries in y locations; the  
105 difference between the coordinates can also serve as normalization to some degree. We also tried  
106 normalizing the data through subtracting the first time point of every vehicle from the entire trajectory.  
107 Therefore, every vehicle will have a similar startpoint of  $[0, 0]$  when feed into the model.

### 108 **2.3.3 Did you use the lane information provided in the dataset. If yes, how did you exploit 109 this information.**

110 No, lane information was not included in our model.

## 111 **3 Deep Learning Model**

### 112 **3.1 Describe the deep learning pipeline for your prediction task and answer the following 113 questions.**

#### 114 **3.1.1 What are the input/output that you end up using for prediction after pre-processing?**

115 Similar to our baseline model, we are choosing to only use positional data when training the deep  
116 learning model. Our output would also only contain positional information. We merge the car and  
117 positional information into one feature vector with the length  $60 * 2 = 120$  as input. The input vector  
118 would also have temporal information which is 19 for input and 30 for output.

#### 119 **3.1.2 What is your loss function? If you have multiple alternatives, discuss your ideas and 120 observations.**

121 The loss function we utilize is the MSE loss because the problem predicts continuous outcome (not a  
122 classification problem) and since trained models have no outlier, MSE can put a larger weight on the  
123 difference between the current location and the previous location.

#### 124 **3.1.3 How did you decide on the deep learning model? Explain the rationale given the 125 input/output.**

126 For the baseline model, we put all the features into a huge array, which is learned by a linear model.  
127 The simple baseline model that we designed is a linear regression model. We are inspired by the idea  
128 of auto regression and teacher-forcing: since we are predicting multiple time points in the future,  
129 each prediction will have a relationship with a previous prediction. We decide to create  $n * 2$  Linear  
130 Regression models,  $n$  = number of time points that need to be predicted. Since we are predicting the  
131 next 3 seconds (30 time points), and we have both x and y coordinates to predict, we would create 60  
132 linear regression models.

133  
134 How we train and test the models are the following: for the x output predictions, we used  
135 all the 30 inputs of x coordinates as features, the first output of x coordinates in the ground truth in  
136 the train data as the label, and we generate parameters  $W$  for the first 19 x inputs to predict the first  
137 output in the test dataset. In the training set, we can always access the ground truth output of both x  
138 and y, but in the testing set, we do not have the output, thus we will use the prediction of the previous  
139 linear regression models to append to the feature matrix to create a new feature matrix with more input.

140  
141 For testing the performance of different design of linear regression models, we used 5942  
142 scenes as validation out of our 205942 data, and we used the validation set to show that the accuracy  
143 dramatically increased when we use the idea of auto regression: if we just predict each output point

144 from the 19 inputs points, the accuracy for the later time points decreased, and the result is much  
145 worse than the autoregression result.

146  
147 While considering which recurrent model would perform best in this task, we first thought  
148 about using a traditional RNN model, possibly LSTM model. Since the data is temporal, the model  
149 would capture this feature with great flexibility unlike a CNN. However, after testing with a basic  
150 LSTM model, we believe that a basic LSTM model is not too good at predicting timesteps when  
151 the predicted timesteps are longer than the inputted timesteps. We seek out to read more papers  
152 about variation of RNN, and we discover that autoencoders/ seq2seq models might be very useful in  
153 our task. Specifically, when using the seq2seq model, there are two LSTM models that are linked  
154 together with a hidden state. More importantly, the decoders can be trained with the teacher forcing  
155 method where existing output labels can sometimes substitute model output, thus increasing the  
156 training accuracy of the model. This type of model fits our tasks well, and we eventually decide on  
157 using the seq2seq model.

158 **3.2 Describe all the models you have tried to make predictions. You should always start**  
159 **with simple models (such as Linear Regression) and gradually increase the complexity**  
160 **of your model.**

161 **3.2.1 Use an itemized list to briefly summarize each of the models, their architecture, and**  
162 **parameters, and provide the correct reference if possible.**

163 1. Close-form Linear Regression

- 164 (a) architecture:  $Y = w * X + b$ 
  - 165 i. Created 60 models, each with different dimensions for X, and w.
- 166 (b) Total number of parameters: 3990
  - 167 i. First model have 36 input columns, with a biased variable, total parameter is 37
  - 168 ii. Total =  $(36 + 1) + (37 + 1) + \dots + (36 + 60 + 1)$

169 2. Multi-layer linear model

- 170 (a) Architecture: (input size: n)  $n = \text{range}(38, 98)$ 
  - 171 i. Linear: (size: n,  $n/2$ )
  - 172 ii. ReLu
  - 173 iii. Linear (size:  $n/2$ ,  $(n/2)/3$ )
  - 174 iv. Linear (size:  $(n/2)/3$ , 1)
  - 175 v. Sigmoid activation function
- 176 (b) Parameters:
  - 177 i. First model,  $n = 36$ , num\_p = 787
  - 178 ii. Last model,  $n = 60$ , num\_p = 5248
  - 179 iii. 161945 in total for all our 60 models

180 3. LSTM model

- 181 (a) Architecture:
  - 182 i. LSTM (4 layers, 520 hidden units, 0.2 dropout rate)
  - 183 ii. Dropout (0.2)
  - 184 iii. Linear (520 -> 120)
- 185 (b) Parameters: 7899960
- 186 (c) Reference: CSE 151B RNN tutorial

187 4. Seq2seq model

- 188 (a) Architecture
  - 189 i. Encoder
    - 190 A. LSTM ( 3 layers, 520 hidden units, 0.2 dropout)
  - 191 ii. Decoder
    - 192 A. Adopt hidden state from the last state of encoder model
    - 193 B. LSTM ( 3 layers, 520 hidden units, 0.2 dropout)
    - 194 C. Linear (520 hidden units -> 120)



195 (b) Parameters  
 196 i. First model have 36 input columns, with a biased variable, total parameter is 37  
 197 ii. Total =  $(36 + 1) + (37 + 1) + \dots + (36 + 60 + 1)$   
 198 (c) Reference: [https://github.com/lkulowski/LSTM\\_encoder\\_decoder](https://github.com/lkulowski/LSTM_encoder_decoder)

199 **3.2.2 If you end up designing your own model architecture, include a picture/sketch of your**  
 200 **model architecture. Explain why you choose such a model.**

201 We used linear regression, multiple-linear model, LSTM, encoder-decoder that are common in Deep  
 202 Learning, therefore, we do not have unique architectures.

203 **3.2.3 Describe different regularization techniques that you have used such as dropout and**  
 204 **max-pooling in your model.**

205 1. Dropout layers are utilized within LSTM and between LSTM and linear mapping layer  
 206 2. Learning rate decay of  $1e-4$  is used while training the model

207 **4 Experiment Design. Describe how you set up the training and testing**  
 208 **design for deep learning.**

209 **4.1 What computational platform/GPU did you use for training/testing?**

210 Schools DSMLP server (datahub), and personal PC setup. We have 1 GPU, 8 CPU, and 16 G RAM.  
 211 For our Local setup: 1 GPU, 32 Gb RAM

212 **4.2 How did you split your training and validation set?**

213 We perform a train/validation split in our train data size of 80% and 20% because the given validation  
 214 set has no label. Thus the train set dimensions is  $164754 * 60 * 19 * 2$ , the validation set dimension is  
 215  $41188 * 60 * 19 * 2$ , and the test set dimension is  $3200 * 60 * 19 * 2$  if we account for all the agents  
 216 19 in the scenes. After we have a relatively good result, we use all train data as the training set and  
 217 test data as the validation set.

218 **4.3 What is your optimizer? How did you tune your learning rate, learning rate decay,**  
 219 **momentum and other parameters?**

220 We tried Adam, SGD, RMSprop optimizers and Adam works the best. We usually set our learning  
 221 rate into 0.01 on linear models and 0.001 on RNN models. In order to tune our learning rate, we used  
 222 a part of the training dataset as a validation set. We also set our learning rate decay as 0.0005. We  
 223 wrap the ADAM optimizer in a scheduler `torch.optim.lr_scheduler.StepLR` and set the parameter step  
 224 size to be 3. Moreover, we used different activation functions to see which can obtain the lower loss.  
 225 We also tried incorporating validation sets into our training, actively tuning the learning rate while  
 226 training with the validation set.

227 **4.4 How did you make multistep (30 step) prediction for each target agent?**

228 We used an RNN approach. Taking advantage of LSTM, we are able to use the information  
 229 of the previous layer to further predict the later frame information for 30 timesteps while  
 230 updating the hidden states H and current state C. Then we add the predicted output to the end  
 231 of the input strain and select the last input-sized timestamps for the next iteration of the model.  
 232 This process iterates until a desirable number of outputs are generated (30 in this case) from the model.  
 233  
 234 Moreover, we also use linear regression/multilayer models through autoregression. The pre-  
 235 diction workflows on the test set goes like this: we used the 19 input x, 19 input y, in total 38 features  
 236 to predict the first x output, using the first linear regression classifier we have made from the training  
 237 set. Then, we used 20 input x (19 input from the real input, and 1 output from the prediction), 19  
 238 input y, in total 39 features to predict the first y output. Then, we used 20 input x, and 20 input y, (40  
 239 features) to predict the second x output.

240 **4.5 How many epoch did you use? What is your batch-size? How long does it take to train**  
241 **your model for one epoch (going through the entire training data set once)?**

242 We usually train 20+ epochs if we just take in a small portion of data. For the Multiple linear model,  
243 our batch size is 512. As we go through the entire training data set, we train 5 epochs for around  
244 80mins (16min/epochs) for linear models. We know that 5 epochs isn't ideal for obtaining the best  
245 accuracy, but due to the time limitations, and the fact that the linear model loss is quite small at the  
246 second epoch, given the massive amount of training data/batches, we used 5 (if number of epoch  
247 increased to 20, we would need many more hours of training). For RNN models, the basic LSTM  
248 model takes around 5 minutes per epoch and the seq2seq model takes around 15 minutes to train one  
249 epoch. For RNN models, batch size is around 256. Basic LSTM model takes around 40 epochs to  
250 convert while autoencoder takes around 10.

251 **4.6 Explain why you made these design choices. Was it motivated by your past experience?**  
252 **Or was it due to the limitation from your computational platform? You are welcome to**  
253 **use screenshots or provide code snippets to explain your design.**

254 It is because of the limitation from the computational platform. The datahub is unstable and always  
255 crashes. Therefore, we can only use the local CPU instead of cuda GPU to train and examine the  
256 model. Most of the design choices come from the compromise among time, computational power, and  
257 performance. If we increase batch size, even though it would decrease time and increase performance,  
258 oftentimes the GPU capacity does not allow. Therefore, we always try to check the maximum batch  
259 size we could run, and then run with a large batch size. Regarding epochs, we observe how the  
260 training loss decreases, and oftentimes we stop the training when the loss stops decreasing and starts  
261 fluctuating.

262 **5 Experiment Results**

263 **5.1 Select a few representative models of yours and report the following results by**  
264 **comparing different models.**

265 **5.1.1 Use a table to compare the performances of different model designs. What conclusions**  
266 **can you draw from this table.**

267

Model	Linear Regression (Closed form)	Multiple Linear Model	Basic LSTM	Autoencoder
Loss	2.53127	10.06080	880.10048	3.67532

268

269 Conclusion: our close form linear regression achieves the best result, this not only shows that a  
270 simple model can achieve the same accuracy as fancier, more complicated models, but also the fact  
271 that our deep learning model needs better design.

272 **5.1.2 Provide an estimate of training time (in flops or minutes) for different models. What**  
273 **did you do to improve the training speed?**

274 For the baseline model, we eliminated the zeros in the dataset to reduce the number of training  
275 objects in order to improve the training speed.

276  
277 For Multi-layer linear, it took 90 mins to go through the entire training data set and train,  
278 and we change the number of layer and activation function to improve the training speed.

279  
280 For the RNN models, different models have different training times according to model  
281 complexity. Regarding the basic LSTM model, one epoch would take around 300 seconds, and the  
282 entire training takes around 40 epochs to reach optimum.

283  
284 For the seq2seq autoencoder model, one w. Also, we increased the batch size and multithreading  
285 number in the model to increase the training speed.

### 5.1.3 Count and report the number of parameters in different models.

- For the linear regression model, as well as the multi-layer linear models, since we want to be greedy and have as much useful training data at predicting each timestamp, we created 60 models to predict each timestamp (30 for x, 30 for y). Therefore, the total number of parameters we count here is the sum of all the number of parameters for all the models.

- Closed form linear regression parameter count: 3990 (sum of 60 models)

- Multi-layer linear model parameters count: 161945 (sum of 60 models)

- Explanation: In the multilinear model in Pytorch, we have 3 layers. Let us take the first model for example, when we have 18 x diff points and 18 y diff points, 36 feature columns in total: for the first layer, we have  $18 * 38$  hidden layer parameters, along with 18 biased parameters, we have  $648 + 18$  parameters, for the second layer, we have  $6 * 18$  hidden layer parameters, along with 6 biased parameters, we have  $108 + 6$  parameters, for the final third layer, we have  $1 * 6$  hidden layer parameters with 1 biased parameters. In total, we have  $648 + 18 + 108 + 6 + 6 + 1 = 787$  parameters; This number varies because we used mathematical formula to determine the parameters of hidden layer based on the input features (if we have 55 features instead of 38 features, the number of hidden layer will also vary; the last model we generate, where we have 96 input points to predict the last x and y timestamp, we will have 5457 parameters.)

- LSTM total Parameters: 7899960

- Encoder-decoder total Parameters:

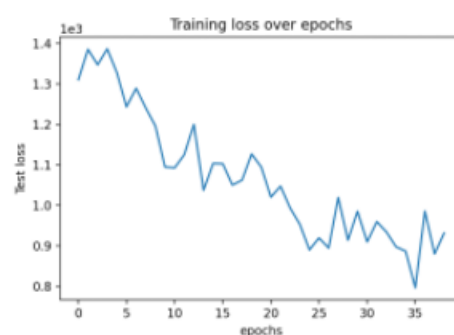
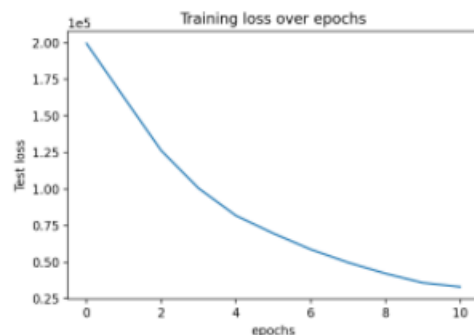
- Encoder: 5500928

- Decoder: 5562488

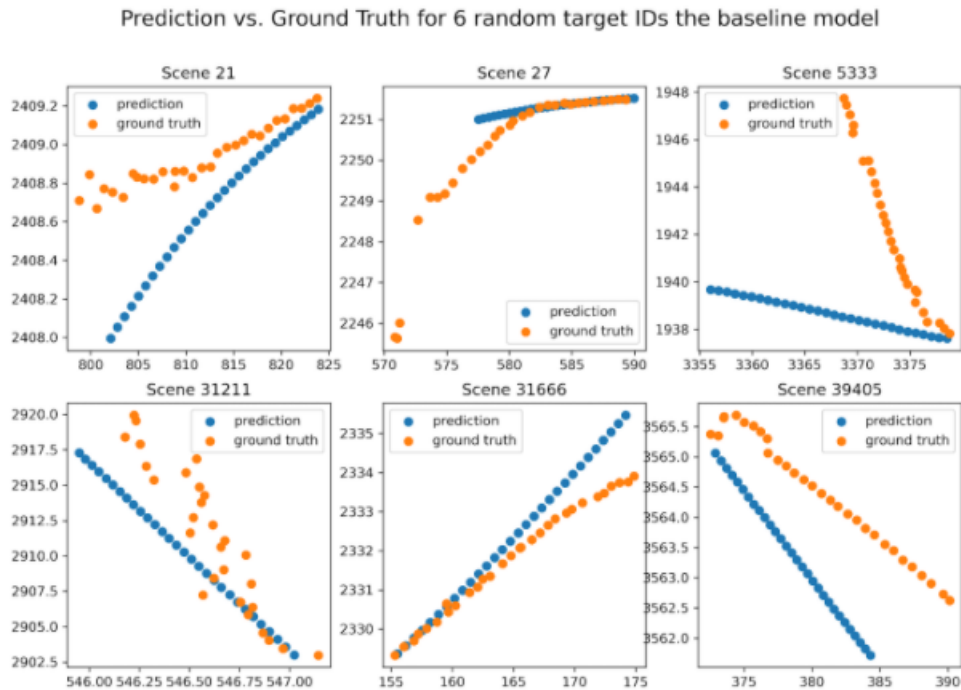
## 5.2 Play with different designs of your model and experiments and report the following for your best-performing design:

### 5.2.1 Visualize the training/validation loss (RMSE) value over training steps (You should expect to see an exponential decay).

The first image is linear regression. The second image is LSTM.



315 **5.2.2 Randomly sample a few training samples after the training has finished. Visualize the**  
 316 **ground truth and your predictions.**



317

318 **5.2.3 Your current ranking on the leaderboard and your final test RMSE.**

319 Private leaderboard: 21 RMSE: 2.67167 Public leaderboard: 20 RMSE: 2.52455

320 **6 Discussion and Future Work :Analyze the results and identify the**  
 321 **lessons/issues that you have learned so far. Briefly answering the following**  
 322 **questions**

323 **6.1 What do you think is the most effective feature engineering strategy?**

324 We use the linear regression model to evaluate which variables/features are important for the prediction  
 325 task. It only takes <15 mins to predict 60 columns , while the DL models take much longer than that.  
 326 After knowing the importance of features, perform data preprocessing and build the deep learning  
 327 model based on the features. In our model, We used this trick to confirm that difference in location  
 328 work better than location, and itself is better than location and velocity combined.

329 **6.2 What techniques (data visualization/model design/hyper-parameter tuning) did you find**  
 330 **most helpful in improving your score?**

331 We think all techniques(data visualization /model design/ hyper-parameter tuning) are all important.  
 332 Data visualization allows us to gain a better understanding of the dataset and the predicted result  
 333 (Geographical visualization).Model design is also very important. We try different models such as  
 334 muti-layer perceptron, Lstm and autoencoder.After you find a good model, tuning the parameter can  
 335 increase your accuracy a bit. Overall, the most helpful one should be data visualization. You can  
 336 not engineer a model and do data processing properly until you really understand your dataset and  
 337 targeted question.

338 **6.3 What was your biggest bottleneck in this project?**

339 Time-consuming. It takes too much time to train and predict the model, and we will not know the  
340 accuracy of model till train though the entire data set. Moreover, familiarity with various machine  
341 learning models. All of the members are not too familiar with RNN, and a lot of time has been put  
342 into understanding the architecture and debugging dimensional problems.

343 **6.4 How would you advise a deep learning beginner in terms of designing deep learning**  
344 **models for similar prediction tasks.**

345 Data Visualization is Important. Getting to know the data is much more important than the training  
346 model. For our model and data, We used the difference between the current location and the previous  
347 location to train in our model instead of involving the velocity and lane information.

348  
349 Keep Everything Simple. Sometimes complex models would make your prediction less in-  
350 terpretable. We try linear regression with careful data preprocessing, and that is still our best outcome  
351 so far.

352  
353 Always Learn from your mistakes. Adjust your time through the mistakes such as high  
354 loss and long time processing. Your mistake can improve your model in a good manner

355 **6.5 If you had more resources, what other ideas would you like to explore?**

356 We want to explore the situation, which the car changes direction. Sometimes cars change their  
357 direction for the next 30 timestamps but our model failed to predict that situation using the first 19  
358 timestamps.

359  
360 Try attention based seq2seq model, which allow the decoder to focus on certain timestamp  
361 by inspecting the attention distribution.

362  
363 Improve the Multi-layer linear model design by adding more validations and convergence  
364 check (since we are creating 60 models, using same number of epochs for them might not work well)

## 365 7 References:

- 366 1. Zhao, Ling, et al. "T-gcn: A temporal graph convolutional network for traffic prediction."  
367 IEEE Transactions on Intelligent Transportation Systems 21.9 (2019): 3848-3858.
- 368 2. Zhao, Zheng, et al. "LSTM network: a deep learning approach for short-term traffic  
369 forecast." IET Intelligent Transport Systems 11.2 (2017): 68-75
- 370 3. Yu, Bing, Haoteng Yin, and Zhanxing Zhu. "Spatio-temporal graph convolutional networks:  
371 A deep learning framework for traffic forecasting." arXiv preprint arXiv:1709.04875 (2017).
- 372 4. Ma, Xiaolei, et al. "Learning traffic as images: a deep convolutional neural network for  
373 large-scale transportation network speed prediction." Sensors 17.4 (2017): 818.
- 374 5. Cui, Zhiyong, et al. "Deep bidirectional and unidirectional LSTM recurrent neural network  
375 for network-wide traffic speed prediction." arXiv preprint arXiv:1801.02143 (2018).

## 376 8 GitHub Link

377 <https://github.com/jeffrey7377/151b-project-ml-noob>